

Stacked Sparse Denoising Autoencoder

In this assignment, we will create a Stacked Sparse Denoising Autoencoder (SSDA) with the purpose of denoising grayscale images.

Not everything required to finish this assignment is covered in class but you should start working on "Imports", "Load Images", and "Create Patches" subsections that follow.

For this assignment, we recommend working in a [Jupyter Notebook \(http://jupyter.org\)](http://jupyter.org). Jupyter Notebook is a browser based interactive programming environment and supports [several languages \(https://github.com/jupyter/jupyter/wiki/Jupyter-kernels\)](https://github.com/jupyter/jupyter/wiki/Jupyter-kernels).

For this assignment, you will most likely need to use a deep learning framework. Several frameworks exist. [Here are a few of them \(https://en.wikipedia.org/wiki/Comparison_of_deep_learning_software\)](https://en.wikipedia.org/wiki/Comparison_of_deep_learning_software). Based on personal and borrowed experience, I will recommend working with Keras if you have no idea what to choose. Brave or experienced users can feel free to use other frameworks.

I (Priyank) will cover an overview of python, numpy, jupyter, and a deep learning framework sometime during next week.

Keras will not work on NEU cluster, though, and you will have to train on your machine or Maria's desktop. If you don't have an access to yet, now is the best time to request it.

Additionally, we will also need to allocate time slots to each of the students to train their models on desktop should someone chooses to implement a complex model.

This homework is worth 100 points. However, there is no breakdown of points per steps. These steps are minimum that is required to score 100 points. Bonus points for showing interesting results and insights.

Credits for this assignment go to Maria's summer advisee Kaleigh O'hara.

This SSDA model is made up of two autoencoders. Rather than train the entire model at once, we train the two autoencoder layers individually, save their weights, and then use those weights to initialize the stacked autoencoder model.

A single autoencoder model includes an input layer, an encoding layer and a decoding layer. After training the first autoencoder and saving its weights, we run the original clean and noisy patches through the input and encoding layers of our training model. The output patches of this model become the input patches to the second autoencoder's training model.

We train the fully assembled SSDA model by initializing the two encoding layers to the weights saved from our individually trained autoencoder models. We provide the trained SSDA model with noisy test images and evaluate its performance.

Keras Blog - 'Deep autoencoder' used as example

<https://blog.keras.io/building-autoencoders-in-keras.html> (<https://blog.keras.io/building-autoencoders-in-keras.html>)

IMPORTS

```
In [1]: #to resolve mkl bug when importing skimage (Seattle Computer Only)
# import mkl
# mkl.get_max_threads()

from keras.layers import Input, Dense
from keras.models import Model

import numpy as np

from keras.datasets import mnist
import numpy as np
import cv2 as cv2

import os as os #some debugging code
# os.sys.path
```

Using TensorFlow backend.

LOAD IMAGES

For this assignment, we will use images from the Berkeley Segmentation Data Set and Benchmarks 500 (BSDS500). The BSDS500 dataset provides three categories of images: train, val, and test, which we will use for training, validating, and testing the models.

<https://www2.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/resources.html> (<https://www2.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/resources.html>)

To reduce the dimensionality of the data, your program should convert the images to greyscale.

Load the BSDS500 train, validate, and test datasets:

```
In [2]: train_path = "/Users/divyaagarwal/Desktop/Computervision/Assignment2Autoe
train_imgs = [f for f in os.listdir(train_path) if f[-3:] == 'jpg']
train_list = [cv2.imread(os.path.join(train_path, f), cv2.IMREAD_GRAYSCALE)
train = np.array([t for t in train_list])
```

```

In [3]: # TRAIN IMAGES (200 Total)
x_train = []

print train[0].shape

for i in range(len(train)):
    x_train.append(train[i].astype('float32'))

for i in range(len(x_train)):
    print x_train[i]

[[ 80.  138.  143. ..., 130.  144.  147.]
 [ 127.  144.  136. ..., 151.  149.  150.]
 [ 137.  133.  133. ..., 150.  150.  153.]
 ...,
 [ 11.   11.   11. ..., 135.  169.  159.]
 [ 11.   11.   11. ..., 131.  182.  187.]
 [ 11.   11.   11. ..., 140.  179.  136.]]
[[ 101.  145.  151. ..., 137.  134.  127.]
 [ 96.   100.  126. ..., 107.  107.  106.]
 [ 80.   93.  137. ..., 38.   39.   41.]
 ...,
 [ 125.  122.  121. ..., 121.  123.  127.]
 [ 127.  104.  102. ..., 121.  114.  108.]
 [ 103.  112.   97. ..., 116.  119.  117.]]
[[ 42.   41.   41. ..., 86.   96.  141.]
 [ 42.   42.   42. ..., 110.  135.  110.]
 [ 42.   42.   43. ..., 126.   94.   83.]
 ...,
 [ 87.   74.   65. ..., 98.   99.   99.]
 [ 87.   72.   66. ..., 72.   97.  105.]

```

```

In [4]: validate_path = "/Users/divyaagarwal/Desktop/Computervision/Assignment2Autoencoder"
validate_imgs = [f for f in os.listdir(validate_path) if f[-3:] == 'jpg']
validate_list = [cv2.imread(os.path.join(validate_path, f), cv2.IMREAD_GRAYSCALE) for f in validate_imgs]
validate = np.array(validate_list)

```

```

In [5]: # VALIDATE IMAGES (100 Total)
x_validate = []

for i in range(len(validate)):
    x_validate.append(validate[i].astype('float32'))

```

```

In [6]: test_path = "/Users/divyaagarwal/Desktop/Computervision/Assignment2Autoencoder"
test_imgs = [f for f in os.listdir(test_path) if f[-3:] == 'jpg']
test_list = [cv2.imread(os.path.join(test_path, f), cv2.IMREAD_GRAYSCALE) for f in test_imgs]
test = np.array(test_list)

```

```
In [7]: # # TEST IMAGES (200 Total)

x_test = []

for i in range(len(test)):
    x_test.append(test[i].astype('float32'))
```

CREATE PATCHES

For each image in the train, validate, and test datasets, create clean and noisy patches using a window size of 8x8 and a step size of 8. Each noisy patch is created by applying random gaussian noise to its equivalent clean patch. Set the seed variable to allow comparison between program executions.

Hint for creating patches: <http://www.pyimagesearch.com/2015/03/23/sliding-windows-for-object-detection-with-python-and-opencv/> (<http://www.pyimagesearch.com/2015/03/23/sliding-windows-for-object-detection-with-python-and-opencv/>)

Output Format:

- Each pixel must be converted from int to float.
- Each patch must be a flattened numpy array of length 64.
- Clean and Noisy sets of Train patches have shape (480000, 64).
- Clean and Noisy sets of Validate patches have shape (240000, 64).
- Clean and Noisy sets of Test patches have shape (480000, 64).

Now creating Clean and Noisy Training patches of shape (480000, 64).

```
In [8]: windowSize = (8,8)
stepSize = 8
# numPatches = 30*2400
numPatches = 480000
numPixels = 64

train_patch = np.empty(shape=(numPatches, numPixels))

k=0
for i in range(len(x_train)):
    image = x_train[i]
    for y in xrange(0, image.shape[0]-8-1, stepSize):
        for x in xrange(0, image.shape[1]-8-1, stepSize):
            train_patch[k] = image[y:y + windowSize[1], x:x + windowSize[1]]
            k = k+1
```

```
In [9]: print train_patch.shape
(480000, 64)
```

```
In [10]: print train_patch[0]
[ 111.  115.  116.  119.  126.  131.  127.  121.  108.  109.  110.  11
  6.
  130.  139.  135.  126.  118.  116.  117.  123.  133.  137.  133.  12
  6.
  123.  122.  126.  133.  134.  131.  130.  132.  125.  126.  132.  13
  8.
  137.  134.  132.  131.  136.  136.  139.  141.  142.  143.  136.  12
  4.
  144.  140.  141.  142.  143.  146.  142.  131.  147.  140.  140.  14
  2.
  138.  138.  142.  142.]
```

```
In [11]: train_noisy_patch = np.empty(shape=(numPatches, numPixels))

# Noisy Patches
sigma = 25

for i in range(len(train_patch)):
    patch = train_patch[i]
    noise = np.random.random(patch.shape) * sigma
    newPatch = np.add(patch, noise)
    train_noisy_patch[i] = np.clip(newPatch, 0, 255)
```

```
In [12]: print train_noisy_patch[0]
[ 131.63920714  132.41504264  133.67948457  131.26625144  140.37018471
  135.47408713  145.07620543  145.64272548  121.3820875   125.299633
  111.85493016  117.96745162  151.61998836  156.95828459  157.40030936
  129.50498264  138.08058669  140.33729645  138.6142917   132.51315001
  133.49610031  152.4027049   150.36458967  135.43869039  131.66695192
  137.84977743  141.63340962  138.97896745  135.65084801  149.95344208
  147.70855627  140.49755504  139.25529465  144.21616729  144.48890336
  155.28312836  140.510952   155.88989093  139.76908771  140.37330441
  144.44581561  157.65962077  156.3865672   162.75398887  156.26158942
  143.95375069  152.33854792  130.91858037  155.26951817  147.04163542
  148.91061268  165.98627358  144.99882658  152.26958255  142.5032973
  136.72919203  155.33356225  164.35548171  160.19363392  146.50334004
  153.58780958  138.94426201  152.86946939  143.61908463]
```

Now creating Clean and Noisy Validate patches of shape (480000, 64).

```
In [13]: windowSize = (8,8)
        stepSize = 8
        #numPatches = 5*2400
        numPatches = 240000
        numPixels = 64

        validate_patch = np.empty(shape=(numPatches, numPixels))

        k=0
        for i in range(len(x_validate)):
            image = x_validate[i]
            for y in xrange(0, image.shape[0]-8-1, stepSize):
                for x in xrange(0, image.shape[1]-8-1, stepSize):
                    validate_patch[k] = image[y:y + windowSize[1], x:x + windowSi
                    k = k+1
```

```
In [14]: print validate_patch.shape

(240000, 64)
```

```
In [15]: print validate_patch[0]

[ 200.  255.  255.  254.  253.  255.  250.  255.  253.  254.  252.  25
5.
  255.  252.  255.  254.  255.  255.  254.  255.  255.  255.  254.  25
5.
  254.  255.  253.  255.  253.  254.  255.  253.  255.  252.  255.  25
5.
  255.  254.  252.  187.  255.  255.  255.  249.  255.  254.  211.  13
3.
  253.  255.  255.  255.  253.  228.  172.  146.  255.  255.  253.  25
4.
  225.  184.  187.  184.]
```

```
In [16]: validate_noisy_patch = np.empty(shape=(numPatches, numPixels))

        # Noisy Patches
        sigma = 25

        for i in range(len(validate_patch)):
            patch = validate_patch[i]
            noise = np.random.random(patch.shape) * sigma
            newPatch = np.add(patch, noise)
            validate_noisy_patch[i] = np.clip(newPatch, 0, 255)
```

In [17]: **print** validate_noisy_patch[0]

```
[ 209.04771447  255.          255.          255.          255.
 255.
 253.59825248  255.          255.          255.          255.
 255.
 255.          255.          255.          255.          255.
 255.
 255.          255.          255.          255.          255.
 255.
 255.          255.          255.          255.          255.
 255.
 255.          255.          255.          202.95793216  255.
 255.
 255.          255.          255.          255.          211.55171391
 146.8584996   255.          255.          255.          255.
 255.
 244.4638734   190.17535621  155.42413838  255.          255.
 255.
 255.          247.8327585   186.79940819  198.46013997  187.56408489
]
```

Now creating Clean and Noisy Test patches of shape shape (480000, 64).

```
In [18]: windowSize = (8,8)
stepSize = 8
#numPatches = 5*2400
numPatches = 480000
numPixels = 64

test_patch = np.empty(shape=(numPatches, numPixels))

k=0
for i in range(len(x_test)):
    image = x_test[i]
    for y in xrange(0, image.shape[0]-8-1, stepSize):
        for x in xrange(0, image.shape[1]-8-1, stepSize):
            test_patch[k] = image[y:y + windowSize[1], x:x + windowSize[0]
            k = k+1
```

In [19]: **print** test_patch.shape

```
(480000, 64)
```

In [20]: **print** test_patch[0]

```
[ 68.  69.  70.  71.  71.  71.  70.  70.  64.  66.  68.  7
0.
 69.  68.  68.  68.  66.  69.  73.  74.  73.  72.  73.  7
4.
 81.  83.  85.  86.  86.  87.  89.  91. 106. 105. 105. 10
5.
107. 109. 112. 114. 122. 120. 118. 117. 118. 121. 123. 12
4.
123. 122. 120. 119. 118. 118. 118. 118. 126. 126. 125. 12
3.
121. 118. 116. 115.]
```

In [21]: test_noisy_patch = np.empty(shape=(numPatches, numPixels))

```
# Noisy Patches
sigma = 25

for i in range(len(test_patch)):
    patch = test_patch[i]
    noise = np.random.random(patch.shape) * sigma
    newPatch = np.add(patch,noise)
    test_noisy_patch[i] = np.clip(newPatch, 0, 255)
```

In [22]: **print** test_noisy_patch.shape

```
(480000, 64)
```

In [23]: **print** test_noisy_patch[0]

```
[ 69.8825226  88.9454007  71.75974664  95.89443798  73.56097471
 77.30254421  94.45172418  88.63319045  77.49728974  85.03947036
 90.91698649  81.12876598  75.10840943  73.65491074  70.33335455
 81.35020673  82.27549456  93.60554131  88.84674512  94.96799633
 76.87117045  76.89172666  90.91571872  86.15355358 100.653349
 85.05844962 103.19592927  88.8988022  104.03830643 110.2727726
105.67517021  98.78557608 111.49248707 116.75635095 122.72021695
108.68634919 121.40105804 130.93875668 119.66934557 127.9529882
124.97947127 142.68872884 119.65153861 134.40535979 126.06842533
140.57910437 144.21381653 147.41734104 135.39619264 131.13019447
134.6437368  143.82939165 123.22631089 129.01074183 124.50801891
122.00851116 135.38097067 133.31102122 128.69968748 133.44744292
130.24687199 122.3277565  133.17853354 129.80312488]
```


SINGLE AUTOENCODER MODELS

Autoencoder Model 1

Train Model 1 - Fit model with random input weights and save output weights

Model takes flattened 8x8 patches, encodes to size 16x16, and decodes to size 8x8.

Parameters:

- verbose = 0 (to prevent keras bug "model.fit ValueError: I/O operation on closed file")
- optimizer = 'adadelta'
- loss = 'mean_squared_error'
- encoder activation = 'relu'
- decoder activation = 'sigmoid'

```
In [24]: # this is the size of our encoded representations
encoding_dim = 256

# this is our input placeholder
input_img = Input(shape=(64,))
# "encoded" is the encoded representation of the input
encoded_model1 = Dense(encoding_dim, activation='relu', kernel_initializer=
    bias_initializer='zeros')(input_img)
# "decoded" is the lossy reconstruction of the input
decoded_model1 = Dense(64, activation='sigmoid', kernel_initializer='rand
    bias_initializer='zeros')(encoded_model1)

# this model maps an input to its reconstruction
autoencoder_model1 = Model(input_img, decoded_model1)
```

```
In [25]: autoencoder_model1.compile(optimizer='adadelta', loss='mean_squared_error')
```

```
In [26]: train_patch = train_patch / 255.
validate_patch = validate_patch / 255.
test_patch = test_patch / 255.
train_noisy_patch = train_noisy_patch / 255.
test_noisy_patch = test_noisy_patch / 255.
validate_noisy_patch = validate_noisy_patch / 255.
```

```
In [27]: history = autoencoder_model1.fit(train_patch, train_patch,
                                         epochs=35,
                                         batch_size=1000,
                                         shuffle=True,
                                         validation_data=(validate_patch, validate_patch))
```

Train on 480000 samples, validate on 240000 samples

Epoch 1/35

480000/480000 [=====] - 6s - loss: 0.0490 - acc: 0.0186 - val_loss: 0.0250 - val_acc: 0.0259

Epoch 2/35

480000/480000 [=====] - 6s - loss: 0.0118 - acc: 0.0577 - val_loss: 0.0087 - val_acc: 0.0535

Epoch 3/35

480000/480000 [=====] - 5s - loss: 0.0078 - acc: 0.0597 - val_loss: 0.0078 - val_acc: 0.0675

Epoch 4/35

480000/480000 [=====] - 5s - loss: 0.0071 - acc: 0.0679 - val_loss: 0.0071 - val_acc: 0.0738

Epoch 5/35

480000/480000 [=====] - 5s - loss: 0.0064 - acc: 0.0735 - val_loss: 0.0064 - val_acc: 0.0777

Epoch 6/35

480000/480000 [=====] - 5s - loss: 0.0058 - acc: 0.0787 - val_loss: 0.0059 - val_acc: 0.0832

Epoch 7/35

480000/480000 [=====] - 5s - loss: 0.0053 - acc: 0.0840 - val_loss: 0.0055 - val_acc: 0.0899

Epoch 8/35

480000/480000 [=====] - 5s - loss: 0.0050 - acc: 0.0903 - val_loss: 0.0052 - val_acc: 0.0957

Epoch 9/35

480000/480000 [=====] - 5s - loss: 0.0048 - acc: 0.0962 - val_loss: 0.0050 - val_acc: 0.1005

Epoch 10/35

480000/480000 [=====] - 5s - loss: 0.0046 - acc: 0.1021 - val_loss: 0.0049 - val_acc: 0.1082

Epoch 11/35

480000/480000 [=====] - 5s - loss: 0.0045 - acc: 0.1079 - val_loss: 0.0047 - val_acc: 0.1138

Epoch 12/35

480000/480000 [=====] - 5s - loss: 0.0043 - acc: 0.1137 - val_loss: 0.0046 - val_acc: 0.1194

Epoch 13/35

480000/480000 [=====] - 5s - loss: 0.0042 - acc: 0.1192 - val_loss: 0.0045 - val_acc: 0.1263

Epoch 14/35

480000/480000 [=====] - 5s - loss: 0.0041 - acc: 0.1243 - val_loss: 0.0043 - val_acc: 0.1320

Epoch 15/35

480000/480000 [=====] - 5s - loss: 0.0040 - acc: 0.1320 - val_loss: 0.0040 - val_acc: 0.1320

```
cc: 0.1296 - val_loss: 0.0042 - val_acc: 0.1371
Epoch 16/35
480000/480000 [=====] - 5s - loss: 0.0038 - a
cc: 0.1343 - val_loss: 0.0041 - val_acc: 0.1448
Epoch 17/35
480000/480000 [=====] - 5s - loss: 0.0037 - a
cc: 0.1391 - val_loss: 0.0040 - val_acc: 0.1468
Epoch 18/35
480000/480000 [=====] - 6s - loss: 0.0036 - a
cc: 0.1431 - val_loss: 0.0039 - val_acc: 0.1497
Epoch 19/35
480000/480000 [=====] - 7s - loss: 0.0035 - a
cc: 0.1481 - val_loss: 0.0038 - val_acc: 0.1556
Epoch 20/35
480000/480000 [=====] - 5s - loss: 0.0034 - a
cc: 0.1522 - val_loss: 0.0037 - val_acc: 0.1600
Epoch 21/35
480000/480000 [=====] - 5s - loss: 0.0034 - a
cc: 0.1568 - val_loss: 0.0036 - val_acc: 0.1657
Epoch 22/35
480000/480000 [=====] - 5s - loss: 0.0033 - a
cc: 0.1608 - val_loss: 0.0036 - val_acc: 0.1681
Epoch 23/35
480000/480000 [=====] - 7s - loss: 0.0032 - a
cc: 0.1644 - val_loss: 0.0035 - val_acc: 0.1739
Epoch 24/35
480000/480000 [=====] - 6s - loss: 0.0031 - a
cc: 0.1676 - val_loss: 0.0034 - val_acc: 0.1773
Epoch 25/35
480000/480000 [=====] - 5s - loss: 0.0031 - a
cc: 0.1711 - val_loss: 0.0034 - val_acc: 0.1808
Epoch 26/35
480000/480000 [=====] - 5s - loss: 0.0030 - a
cc: 0.1745 - val_loss: 0.0033 - val_acc: 0.1851
Epoch 27/35
480000/480000 [=====] - 5s - loss: 0.0030 - a
cc: 0.1778 - val_loss: 0.0032 - val_acc: 0.1859
Epoch 28/35
480000/480000 [=====] - 5s - loss: 0.0029 - a
cc: 0.1812 - val_loss: 0.0032 - val_acc: 0.1906
Epoch 29/35
480000/480000 [=====] - 5s - loss: 0.0029 - a
cc: 0.1842 - val_loss: 0.0031 - val_acc: 0.1937
Epoch 30/35
480000/480000 [=====] - 5s - loss: 0.0028 - a
cc: 0.1874 - val_loss: 0.0031 - val_acc: 0.1968
Epoch 31/35
480000/480000 [=====] - 5s - loss: 0.0028 - a
cc: 0.1909 - val_loss: 0.0030 - val_acc: 0.2004
```

```

Epoch 32/35
480000/480000 [=====] - 6s - loss: 0.0027 - a
cc: 0.1940 - val_loss: 0.0030 - val_acc: 0.2031
Epoch 33/35
480000/480000 [=====] - 6s - loss: 0.0027 - a
cc: 0.1971 - val_loss: 0.0029 - val_acc: 0.2059
Epoch 34/35
480000/480000 [=====] - 6s - loss: 0.0026 - a
cc: 0.2001 - val_loss: 0.0029 - val_acc: 0.2079
Epoch 35/35
480000/480000 [=====] - 5s - loss: 0.0026 - a
cc: 0.2033 - val_loss: 0.0028 - val_acc: 0.2132

```

```

In [28]: # save the weights
weights_model1 = []
weights_model1.append(autoencoder_model1.get_weights())

print len(weights_model1[0])

```

4

Plot the Loss from Model 1

```

In [29]: import matplotlib.pyplot as plt

print(history.history.keys())

# summary for accuracy



plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

# summary for loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

['acc', 'loss', 'val_acc', 'val_loss']

```

model accuracy



Feedforward Model 1

Feed clean and noisy patches through the input and encoding layers of model 1.

Note: The weights of the encoding layer must equal the encoding layer weights obtained from the trained model 1.

This model will transform your flattened 8x8 images to flattened 16x16 images, which will be used as input to Autoencoder Model 2.

```
In [30]: # this model maps an input to its encoded representation
encoder_model1 = Model(input_img, encoded_model1)

#load weights from previous model
encoder_model1.set_weights([weights_model1[0][0], weights_model1[0][1]])

#compile the model
encoder_model1.compile(optimizer='adadelta', loss='mean_squared_error')

# save the O/P 16x16
encoder_clean = encoder_model1.predict(train_patch)

print encoder_clean.shape

# save the O/P 16x16
encoder_noisy = encoder_model1.predict(train_noisy_patch)

print encoder_noisy.shape

(480000, 256)
(480000, 256)
```

Autoencoder Model 2

Train Model 2 - Fit model with random input weights and save output weights

Model takes flattened 16x16 patches, encodes to size 32x32, and decodes to size 16x16.

Parameters:

- verbose = 0 (to prevent keras bug "model.fit ValueError: I/O operation on closed file")
- optimizer = 'adadelta'
- loss = 'mean_squared_error'
- encoder activation = 'relu'
- decoder activation = 'sigmoid'

```
In [34]: # this is the size of our encoded representations
encoding_dim = 1024

# this is our input placeholder
input_img = Input(shape=(256,))

# "encoded" is the encoded representation of the input
encoded_model2 = Dense(encoding_dim, activation='relu', kernel_initializer='zeros')(input_img)

# "decoded" is the lossy reconstruction of the input
decoded_model2 = Dense(256, activation='sigmoid', kernel_initializer='zeros')(encoded_model2)

# this model maps an input to its reconstruction
autoencoder_model2 = Model(input_img, decoded_model2)
```

```
In [35]: autoencoder_model2.compile(optimizer='adadelta', loss='mean_squared_error')
```

```
In [36]: history = autoencoder_model2.fit(encoder_clean,
                                          encoder_clean,
                                          batch_size=1000,
                                          epochs=40,
                                          validation_split=0.33)
```

Train on 321599 samples, validate on 158401 samples

Epoch 1/40

321599/321599 [=====] - 31s - loss: 0.0535 - val_loss: 0.0338

Epoch 2/40

321599/321599 [=====] - 31s - loss: 0.0386 - val_loss: 0.0288

Epoch 3/40

321599/321599 [=====] - 31s - loss: 0.0374 - val_loss: 0.0276

Epoch 4/40

321599/321599 [=====] - 31s - loss: 0.0369 - val_loss: 0.0269

Epoch 5/40

321599/321599 [=====] - 32s - loss: 0.0365 - val_loss: 0.0265

Epoch 6/40

321599/321599 [=====] - 38s - loss: 0.0362 - val_loss: 0.0261

Epoch 7/40

```
In [37]: autoencoder_model2.save_weights("/Users/divyaagarwal/Desktop/Computervisi")
```

```
In [38]: # save the weights
weights_model2 =[]
weights_model2.append(autoencoder_model2.get_weights())

print len(weights_model2[0])
```

4


```
In [42]: import matplotlib.pyplot as plt

print(history.history.keys())

plt.plot(history.history['loss'])

# summary for loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```

['loss', 'val_loss']



Feedforward Model 2

Feed the clean and noisy patches from the feedforward model 1 through the input and encoding layers of model 2.

Note: The weights of the encoding layer must equal the encoding layer weights obtained from the trained model 2.

This model will transform your flattened 16x16 images to flattened 32x32 images, which can be used for testing and debugging the trained model.

```
In [43]: # this is the size of our encoded representations
encoding_dim = 1024
input_img = Input(shape=(256,))

encoded_model2 = Dense(encoding_dim, activation='relu')(input_img)

# this model maps an input to its encoded representation
encoder_model2 = Model(input_img, encoded_model2)

#set weights
encoder_model2.set_weights([weights_model2[0][0], weights_model2[0][1]])

#compile
encoder_model2.compile(optimizer='adadelta', loss='mean_squared_error')

# save the O/P 32x32
encoder_clean2 = encoder_model2.predict(encoder_clean)

print encoder_clean2.shape

# save the O/P 32x32
encoder_noisy2 = encoder_model2.predict(encoder_noisy)

print encoder_noisy2.shape

(480000, 1024)
(480000, 1024)
```

STACKED SPARSE DENOISING AUTOENCODER

Train the SSDA

The SSDA model will take patches of size 8x8 as input.

The first encoding layer expands an 8x8 patch to 16x16. The second encoding layer expands a 16x16 patch to 32x32. We then introduce the two decoding layers, which decrease the image from 32x32 to 16x16 to 8x8. The purpose of our model is take a noisy 8x8 image as input and return a clean 8x8 image as output.

The weights of the two encoding layers must be initialized with the encoding layer weights that were saved from the corresponding single autoencoder models above.

```
In [44]: input_img1 = Input(shape=(64,))

encoded1 = Dense(256, activation='relu')(input_img1)

encoded2 = Dense(1024, activation='relu')(encoded1)

# "decoded" used encoded i/p of noisy data 32x32 dim
decoded1 = Dense(256, activation='sigmoid', kernel_initializer='random_no
            bias_initializer='zeros')(encoded2)

decoded2 = Dense(64, activation='sigmoid', kernel_initializer='random_nor
            bias_initializer='zeros')(decoded1)
```

```
In [46]: # send the model
autoencoder_ssda = Model(input_img1, decoded2)

# set the weights 1
autoencoder_ssda.layers[1].set_weights([weights_model1[0][0], weights_mod
# set the weights 2
autoencoder_ssda.layers[2].set_weights([weights_model2[0][0], weights_mod

# compile it now
autoencoder_ssda.compile(optimizer='adadelata', loss='mean_squared_error',

# weights_ssda =[]
# weights_ssda.append(autoencoder_ssda.get_weights())

# autoencoder_ssda.set_weights([weights_model1[0][0],
#                               weights_model1[0][1],
#                               weights_model2[0][0],
#                               weights_model2[0][1],
#                               weights_ssda[0][4],
#                               weights_ssda[0][5],
#                               weights_ssda[0][6],
#                               weights_ssda[0][7]])
```

```
In [47]: history = autoencoder_ssda.fit(train_noisy_patch,
                                         train_patch,
                                         batch_size=1000,
                                         epochs=40,
                                         validation_data=(validate_noisy_patch, val
```

Train on 480000 samples, validate on 240000 samples

Epoch 1/40

480000/480000 [=====] - 65s - loss: 0.0165 -
acc: 0.0231 - val_loss: 0.0094 - val_acc: 0.0339

Epoch 2/40

480000/480000 [=====] - 54s - loss: 0.0080 -
acc: 0.0537 - val_loss: 0.0081 - val_acc: 0.0521

Epoch 3/40

480000/480000 [=====] - 57s - loss: 0.0069 -
acc: 0.0510 - val_loss: 0.0069 - val_acc: 0.0517

Epoch 4/40

480000/480000 [=====] - 77s - loss: 0.0059 -
acc: 0.0533 - val_loss: 0.0061 - val_acc: 0.0559

Epoch 5/40

480000/480000 [=====] - 57s - loss: 0.0054 -
acc: 0.0576 - val_loss: 0.0057 - val_acc: 0.0590

Epoch 6/40

480000/480000 [=====] - 55s - loss: 0.0051 -
acc: 0.0617 - val_loss: 0.0054 - val_acc: 0.0638

Epoch 7/40

Save SSDA Model Weights

```
In [49]: autoencoder_ssda.save_weights("/Users/divyaagarwal/Desktop/Computervision
```

Plot the Loss from SSDA

```
In [50]: import matplotlib.pyplot as plt



print(history.history.keys())

# summarize history for accuracy

plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```

```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```

```
['acc', 'loss', 'val_acc', 'val_loss']
```



Run SSDA on Noisy Test Images and Calculate Mean Squared Error

```
In [51]: final_prediction = autoencoder_ssda.predict(test_noisy_patch)
```

```
In [52]: autoencoder_ssda.save_weights("/Users/divyaagarwal/Desktop/Computervision
#mean_squared_error(test_patch, final_prediction)
```

VISUALIZE RESULTS

Use the image 119082.jpg from the validate image set to visualize effectiveness of the SSDA model.

Load and Display the image

```
In [53]: import matplotlib.pyplot as plt
img = validate[10]
print img
plt.imshow(img, cmap='gray')
plt.axis('off')
plt.show()
```

```
[[ 38  37  37 ..., 171 170 170]
 [ 31  30  30 ..., 169 170 171]
 [ 29  26  25 ..., 171 170 170]
 ...,
 [255 243 231 ..., 248 253 246]
 [125  71  41 ..., 252 255 246]
 [ 66  77  55 ..., 206 182 114]]
```



Create patches from the image

Select a patch with significant contrast (for example, a clearly defined vertical or horizontal line).


```
In [54]: clean_patch_img = validate_patch[2400*9 : 2400*10 -1]
print clean_patch_img[0]
```


```
[ 0.23529412  0.23921569  0.24313725  0.24313725  0.23921569  0.235294
12
 0.22745098  0.22745098  0.23921569  0.23921569  0.23921569  0.239215
69
 0.23529412  0.23137255  0.22352941  0.21960784  0.24313725  0.243137
25
 0.24313725  0.23921569  0.23529412  0.22745098  0.22352941  0.219607
84
 0.24313725  0.24705882  0.24313725  0.23921569  0.23137255  0.227450
98
 0.22352941  0.22352941  0.23921569  0.24313725  0.24705882  0.239215
69
 0.22352941  0.21568627  0.21960784  0.22352941  0.24313725  0.247058
82
 0.24705882  0.23921569  0.22352941  0.21568627  0.21568627  0.223529
41
 0.24313725  0.24705882  0.24705882  0.23921569  0.22745098  0.219607
84
 0.21960784  0.22352941  0.23921569  0.23921569  0.23921569  0.235294
12
 0.22745098  0.22352941  0.21960784  0.21960784]
```

```
In [55]: noisy_patch_img = validate_noisy_patch[2400*9 : 2400*10 -1]
```

Display the clean patch




```
In [56]: for i in range(4):
    patch = clean_patch_img[i]
    patch = patch.reshape(8,8) * 255.0
    plt.imshow(patch, cmap='gray')
    plt.axis('off')
    plt.show()
```







Display the noisy patch

```
In [57]: for i in range(4):  
    patch = noisy_patch_img[i]  
    patch = patch.reshape(8,8) * 255.0  
    plt.imshow(patch, cmap='gray')  
    plt.axis('off')  
    plt.show()
```









Feed the noisy patch through the trained SSDA model and display the resulting denoised image

Does the denoised patch resemble the clean patch?

```
In [58]: output = final_prediction[2400*9 : 2400*10-1]
```

```
In [59]: for i in range(4):
          patch = output[i]
          patch = patch.reshape(8,8) * 255.0
          plt.imshow(patch, cmap='gray')
          plt.axis('off')
          plt.show()
```





Plot a single image patch and compare

ANALYZE RESULTS

Comment on the Effectiveness of your model.

Can you improve the model by changing any of the parameters?

```
In [68]: print output[1234].reshape(8,8) * 255

[[ 37.47543716  37.89933395  36.84737015  34.51482773  35.21337891
   33.31004715  31.19306374  32.2495079 ]
 [ 37.3486557   37.1876297   35.6241684   35.13164139  34.55157852
   33.40540314  31.43386269  31.19120026]
 [ 35.12536621  34.93384552  34.81932831  34.60026932  33.70185471
   33.41449738  32.11282349  32.08383179]
 [ 33.25758743  32.35405731  32.54502869  33.20717621  33.28340149
   34.02289963  32.88191223  31.97543907]
 [ 29.74655914  29.48270226  29.38748741  29.52987671  31.82883072
   32.26580048  32.11250687  31.10915756]
 [ 29.54504776  28.71029854  28.3489666   28.59669495  29.91965485
   30.5398941   31.06070328  30.09921074]
 [ 28.55385017  27.87735748  28.20727158  27.90884018  29.74308014
   30.32805252  30.73750305  30.86085129]
 [ 31.33398628  29.25794983  29.29073715  28.75881386  29.35293961
   30.19287682  30.62119484  32.3893013 ]]
```

```
In [81]: #add vmax and vmin : todo

plt.figure(1)
ax = plt.subplot(121)
patch_clean = clean_patch_img[1234].reshape(8,8) * 255.0
plt.imshow(patch_clean, cmap='gray', vmin=25, vmax=160)
ax.set_title('Clean')
plt.axis('off')
plt.colorbar()

ax = plt.subplot(122)
patch_noisy = noisy_patch_img[1234].reshape(8,8) * 255.0
plt.imshow(patch_noisy, cmap='gray', vmin=25, vmax=160)
ax.set_title('Noisy')
plt.axis('off')
plt.colorbar()
plt.show()

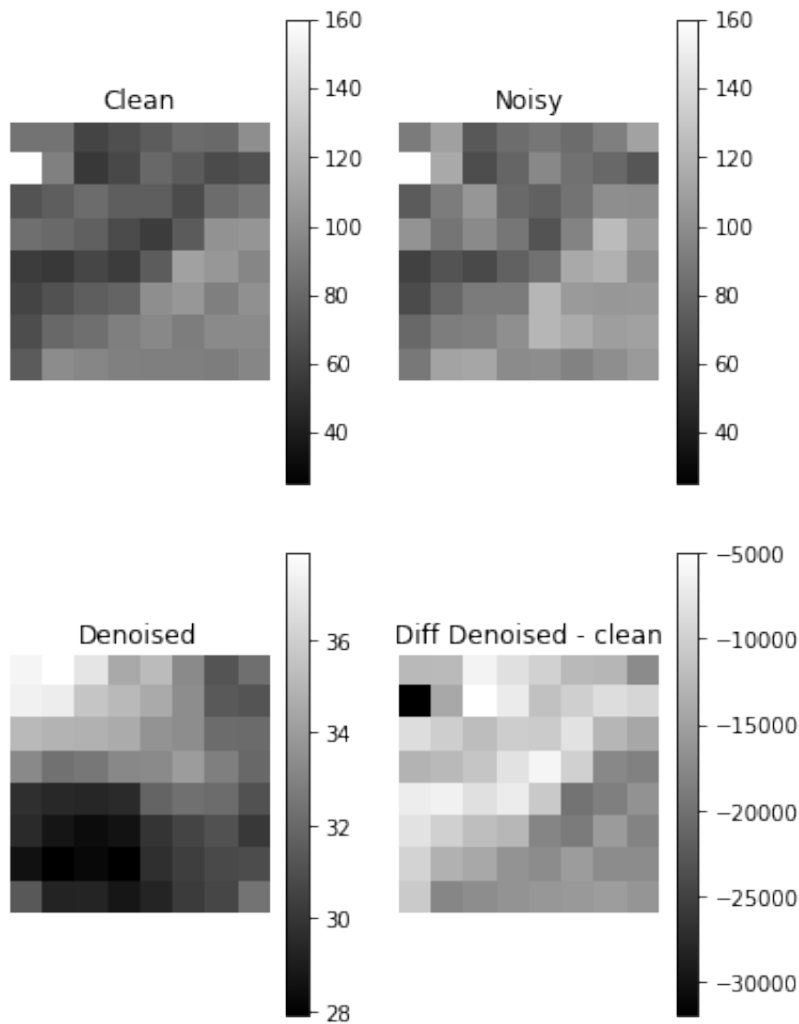
ax = plt.subplot(121)
```

```

patch_denoised = output[1234].reshape(8,8) * 255.0
plt.imshow(patch_denoised, cmap='gray')
ax.set_title('Denoised')
plt.axis('off')
plt.colorbar()

ax =plt.subplot(122)
patch_noisy_clean_diff = patch_denoised - patch_clean
patch_noisy_clean_diff = patch_noisy_clean_diff.reshape(8,8) * 255.0
plt.imshow(patch_noisy_clean_diff, cmap='gray') #vmin=-10000, vmax=-60000
ax.set_title('Diff Denoised - clean')
plt.colorbar()
plt.axis('off')
plt.show()

```



```

In [73]: #Calculate psnr 1
psnr1 = -10. * np.log10(np.mean(np.square(patch_denoised - patch_clean)))

```

```
In [74]: #calculate psnr 2
psnr2 = -10. * np.log10(np.mean(np.square(patch_noisy - patch_clean)))
```

```
In [75]: print "patch_denoised - patch_clean"
print psnr1
print "patch_noisy - patch_clean"
print psnr2
```

```
patch_denoised - patch_clean
-34.5942049317
patch_noisy - patch_clean
-22.7518845712
```

The model could be more effective if it could be trained on higher epoch.

By increasing the epoch the efficiency can be increased. We might we overfitting with the small dataset, we could improve and avoid overfitting it by using a bigger datasets. Also using a bigger CNN might increase the accuracy and decrease the loss.